

Vaibhav Shekhar Dave  
Large Scale Parallel Data Processing  
HW-2 Report

Git repo:

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/tree/master/SocialTraingle>

## Implementation:

### Psuedo Code:

1) Reduce-side Join with Max filter -

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/blob/master/SocialTraingle/MR-Demo/src/main/java/wc/TriangleCount.java>

Let the data set be key value pair for(userID, followerID) represenated by each row.

-----  
Map1 for finding Single Path

```
if(userID) < maxFilter && follower < maxFilter{  
    {  
    emit( userID, "userid" - "follower"- "from")  
    emit( followerID, "userid" - "follower"- "to")  
    }
```

-----  
Reducer1 for Path2 – Partitioned by keys

Collect all to-s and from in two different lists namely to and from.

As the key is same we the follower info for a particular user in from list  
and the information about whom a particular user follows in to List.

```
for (Text fromEdge : from) {  
    for (Text toEdge : to) {  
        // from is a->b and toEdge is b->c  
        // So we want to emit a->c  
  
        context.write(new Text(toEdge.toString().split("-")[0]), // a  
            new Text(fromEdge.toString().split("-")[1])); // c  
    }  
}
```

---

Map 2.1 – reiterates through the database to make a to List

```
if(userID) < maxFilter && follower < maxFilter{  
  for each record emit (follower+"-"+userID, "to")  
}
```

---

Map 2.2 – iterates through the reduce1 output to make a Path2List List

Let the output of reduce1 be in the format (a,b)

for each row

```
emit(a+"-"+b, from)
```

Final Reducer: Partitioned by keys

Takes input from output of Map2.1 and Map2.2

For every row

Calculate number of Tos and number of froms

So To.count > 0

There is path between that means we have completed the triangle  
and add the number of fromCounts to the main count (global variable).

---

## 2) Map-Side Join:

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/blob/master/SocialTriangle/MR-Demo/src/main/java/wc/TwitterRepJoin.java>

Map 1: Let the data set be key value pair for(userID, followerID) represented by each row.

Setup() -

Iterate through all rows and make a HashMap <userID, List<FollowerID>>

For max filter

```
if(userID > max || followerID > max) {  
    // do nothing  
}  
else{  
    map.add(userID, List.add(FollowerID));  
}
```

### Map

Read each row from input: (userID, followerID)

```
if(userID > max || followerID > max) {  
    // do nothing  
}
```

for each userID1 in row find userID1.followers(from map)

for each userID2 in userID1.followers find userID2.followers(from map)

```
if(userID2.followers.contains(userID)){  
    count++  
}
```

Configuration :

Small Cluster (4 nodes) :

REP-join, MAX = 20000 - Number of triangles 1315197755 time = 58 min

RS-join, MAX = 10000 - Number of triangles 7234833 time = 15 min

Large Cluster (7 nodes) :

REP-join, MAX = 20000 - Number of triangles 1315197755 time = 33 min

RS-join, MAX = 10000 - Number of triangles 7234833 time = 8 min

Output from Twitter Replicated Join:

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/tree/d2dcb9748b38515cc183d80d00b00bfa47ca60f5/SocialTraingle/MR-Demo/output>

Output from RS Join:

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/tree/d375f9319bb62890a3ee8785998a37d38881022d/SocialTraingle/MR-Demo/output-Triangle>

## Logs

### Twitter Rep Join – Large Cluster

[https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2UTY17W9WITC2/hadoop-mapreduce/history/2019/02/10/000000/job\\_1549838415613\\_0001-1549838663252-hadoop-TwitterRepJoin-1549840651485-20-0-SUCCEEDED-default-1549838710824.jhist.gz](https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2UTY17W9WITC2/hadoop-mapreduce/history/2019/02/10/000000/job_1549838415613_0001-1549838663252-hadoop-TwitterRepJoin-1549840651485-20-0-SUCCEEDED-default-1549838710824.jhist.gz)

### RSJoin job 1 - Large Cluster

[https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2UTY17W9WITC2/hadoop-mapreduce/history/2019/02/10/000000/job\\_1549838415613\\_0002-1549841029486-hadoop-RSJoin-1549841186902-20-11-SUCCEEDED-default-1549841035875.jhist.gz](https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2UTY17W9WITC2/hadoop-mapreduce/history/2019/02/10/000000/job_1549838415613_0002-1549841029486-hadoop-RSJoin-1549841186902-20-11-SUCCEEDED-default-1549841035875.jhist.gz)

### RS-Join Complete - Large Cluster

[https://s3.console.aws.amazon.com/s3/object/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2UTY17W9WITC2/hadoop-mapreduce/history/2019/02/10/000000/job\\_1549838415613\\_0003-1549841188946-hadoop-RSJoin%252BComplete%252BTriangle-1549841504069-47-11-SUCCEEDED-default-1549841195587.jhist.gz?region=us-east-1&tab=overview](https://s3.console.aws.amazon.com/s3/object/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2UTY17W9WITC2/hadoop-mapreduce/history/2019/02/10/000000/job_1549838415613_0003-1549841188946-hadoop-RSJoin%252BComplete%252BTriangle-1549841504069-47-11-SUCCEEDED-default-1549841195587.jhist.gz?region=us-east-1&tab=overview)

### Twitter Rep join

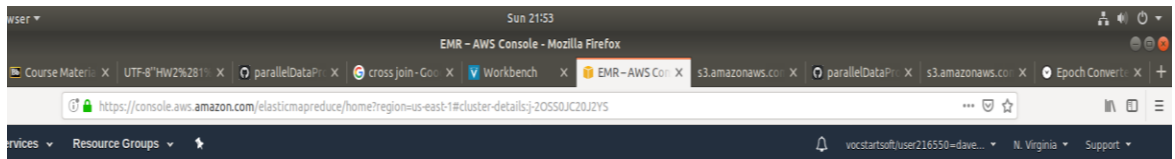
[https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2OSS0JC20J2YS/hadoop-mapreduce/history/2019/02/10/000000/job\\_1549811928896\\_0011-1549818237189-hadoop-TwitterRepJoin-1549821689429-20-0-SUCCEEDED-default-1549818243398.jhist.gz](https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2OSS0JC20J2YS/hadoop-mapreduce/history/2019/02/10/000000/job_1549811928896_0011-1549818237189-hadoop-TwitterRepJoin-1549821689429-20-0-SUCCEEDED-default-1549818243398.jhist.gz)

### RS Join 1

[https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2OSS0JC20J2YS/hadoop-mapreduce/history/2019/02/10/000000/job\\_1549811928896\\_0002-1549812171266-hadoop-RSJoin-1549812429287-20-5-SUCCEEDED-default-1549812179685.jhist.gz](https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2OSS0JC20J2YS/hadoop-mapreduce/history/2019/02/10/000000/job_1549811928896_0002-1549812171266-hadoop-RSJoin-1549812429287-20-5-SUCCEEDED-default-1549812179685.jhist.gz)

### RS join Complete:

[https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2OSS0JC20J2YS/hadoop-mapreduce/history/2019/02/10/000000/job\\_1549811928896\\_0003-1549812432231-hadoop-RSJoin%2BComplete%2BTriangle-1549813064651-44-5-SUCCEEDED-default-1549812438345.jhist.gz](https://s3.amazonaws.com/aws-logs-577453344208-us-east-1/elasticmapreduce/j-2OSS0JC20J2YS/hadoop-mapreduce/history/2019/02/10/000000/job_1549811928896_0003-1549812432231-hadoop-RSJoin%2BComplete%2BTriangle-1549813064651-44-5-SUCCEEDED-default-1549812438345.jhist.gz)



Cluster: hadoopAssignment2-1 Terminated Terminated by user request

Summary Application history Monitoring Hardware Events Steps Configurations Bootstrap actions

Amazon EMR collects information from YARN applications on your cluster and keeps historical information for up to seven days after applications have completed. Detailed application history is only available for Spark. [Learn more](#)

YARN applications (7)

Filter: All applications Filter applications... 7 applications (all loaded)

Application ID	Type	Action	Status	Start time (UTC-5)	Duration	Finish time (UTC-5)	User
application_1549811928896_0011	MapReduce	TwitterRepJoin	Succeeded	2019-02-10 12:03 (UTC-5)	58 min	2019-02-10 13:01 (UTC-5)	hadoop
application_1549811928896_0010	MapReduce	TwitterRepJoin	Succeeded	2019-02-10 11:52 (UTC-5)	4.4 min	2019-02-10 11:56 (UTC-5)	hadoop
application_1549811928896_0009	MapReduce	TwitterRepJoin	Failed	2019-02-10 11:15 (UTC-5)	1.3 min	2019-02-10 11:17 (UTC-5)	hadoop
application_1549811928896_0005	MapReduce	TwitterRepJoin	Failed	2019-02-10 10:58 (UTC-5)	1.4 min	2019-02-10 11:00 (UTC-5)	hadoop
application_1549811928896_0004	MapReduce	TwitterRepJoin	Failed	2019-02-10 10:44 (UTC-5)	1.9 min	2019-02-10 10:45 (UTC-5)	hadoop
application_1549811928896_0003	MapReduce	RSJoin Complete Triangle	Succeeded	2019-02-10 10:27 (UTC-5)	11 min	2019-02-10 10:37 (UTC-5)	hadoop
application_1549811928896_0002	MapReduce	RSJoin	Succeeded	2019-02-10 10:22 (UTC-5)	4.3 min	2019-02-10 10:27 (UTC-5)	hadoop

Cluster: BigCluster3 Terminated Terminated by user request

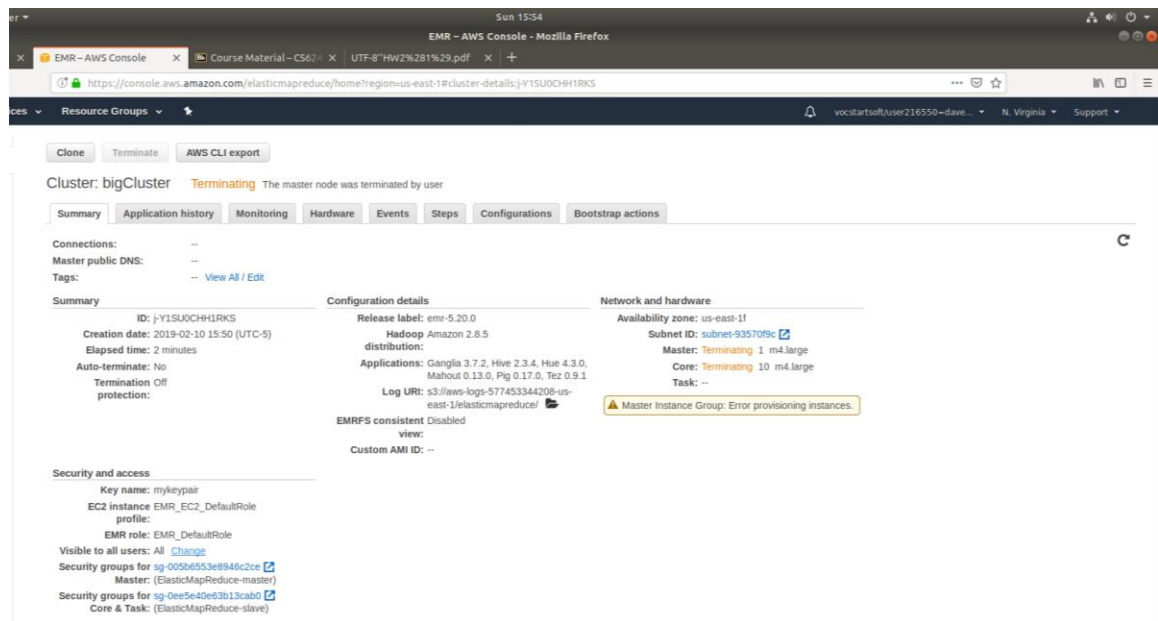
Summary Application history Monitoring Hardware Events Steps Configurations Bootstrap actions

Amazon EMR collects information from YARN applications on your cluster and keeps historical information for up to seven days after applications have completed. Detailed application history is only available for Spark. [Learn more](#)

YARN applications (3)

Filter: All applications Filter applications... 3 applications (all loaded)

Application ID	Type	Action	Status	Start time (UTC-5)	Duration	Finish time (UTC-5)	User
application_1549838415613_0003	MapReduce	RSJoin Complete Triangle	Succeeded	2019-02-10 18:26 (UTC-5)	5.3 min	2019-02-10 18:31 (UTC-5)	hadoop
application_1549838415613_0002	MapReduce	RSJoin	Succeeded	2019-02-10 18:23 (UTC-5)	2.6 min	2019-02-10 18:26 (UTC-5)	hadoop
application_1549838415613_0001	MapReduce	TwitterRepJoin	Succeeded	2019-02-10 17:44 (UTC-5)	33 min	2019-02-10 18:17 (UTC-5)	hadoop



Tried using 11 nodes but I got this error so my small cluster has 4 nodes and large has 7 nodes.

## Analysis:

Show the MapReduce pseudo-code for the program you used to determine the cardinality (and maybe data volume) of Path2. If you did not use a program, show the steps of the analysis you performed to estimate the number?

### 1) Psuedo Code

Mapper

Let each row of the record be (userid , followerID)

Emit(userid, from)

Emit(followid, to)

Reducer

Count the number of from c

Count the number of to d

Mul = c\*d

Total count(global var) += mul

Implementation:

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/blob/master/SocialTraingle/MR-Demo/src/main/java/wc/Path2JoinRSCardinality.java>

2) Show the table with all 12 cardinality and all 12 volume estimates for the two join steps and RS-join vs. Rep-join. If you merge the two steps into one for one or both join types, state so clearly in the report. Then you only have to report the corresponding input/shuffle/file cache/output numbers for the merged program

	RS Join input	RS join Shuffle	RS join Output	Rep Join input	Rep join File Cache	Rep join Output
Step 1	Full input of Edges.csv  1319496574 bytes	1356928 Bytes  88571 records	1302114799 Bytes  4387766 records	1319496574 bytes	9130999808 bytes	1315197755 bytes
Step 2	4387766 records  1302114799 Bytes	26397 records  23200210944	<b>201519176792 Records</b>  <b>Ran on whole edges.csv</b>	Merged step 1 and 2	Merged step 1 and 2	Merged step 1 and 2



3) For cardinality of Path2 I have the following code and output

Output:

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/blob/915bb51456220b3d589d923521becd5b172cf214/SocialTraingle/MR-Demo/output/part-r-00000>

Code:

<https://github.ccs.neu.edu/vaibhavdave5/parallelDataProcessing/blob/915bb51456220b3d589d923521becd5b172cf214/SocialTraingle/MR-Demo/src/main/java/wc/Path2JoinRSCardinality.java>